



Officials cluster with their namesakes, and the rarer the name the more they cluster

Left: how much more surname-concentrated each ministry's staff is than the same people would be if surnames were dealt out at random within entry decade. All 30 ministries sit above chance, 22 of them by more than any of 200 permutations reached; the spread runs from human rights at 3.26× down to interior at 1.05×, which is not separable from chance. Right: how often someone arriving at a body finds a namesake already there — 14.9% of 98,957 arrivals — split by how many people in the register carry the surname. The excess over chance runs from 2.78× for a surname held by one or two people to 1.15× for one held by more than sixty. A shared rare name is evidence of a shared family; a shared common one mostly is not, and that is where the effect goes.

Within a ministry



On arrival at any body



Source: Journal Officiel de la République Tunisienne, 1957-2026 (6,378 French issues, jort.tn). Author's extraction. A surname is read as everything from a particle to the end of the name, so Ben Abdallah is not filed under Abdallah. Which particle takes two rules, because the two kinds behave differently in the record: ben, bin and ould chain, and 1,078 names run 'X ben Y ben Z', where Tunisian usage makes the terminal patronymic the family name, so the last of them opens the surname; el, bel, bou, hadj, caïd and the rest also occur inside compound given names — el is followed by a later ben in 117 names — so one of these opens the surname only when no ben, bin or ould is present, and then it is the first. Taking the first particle regardless files Zine El Abidine Ben Ali, the most-recorded name in the corpus, under 'el abidine ben ali'. Where there is no particle the surname is the final token. People whose final token is a bare given name are left out, because the record does not say whether it is a family name or a second given name. So are the rows where the extractor took French running text from the surrounding act for a name. Posts whose portfolio names no ministry are not a ministry and are not drawn. The proxy is weak in both directions and the weakness is not symmetric in its consequences: 77% of the 11,850 people recorded under more than one spelling have variants that disagree about their own surname, which files real kin under different names; and the commonest surnames here — Trabelsi, Cherif, Dridi, Hammami, Gharbi — are the commonest surnames in Tunisia, which files strangers under one name. Both push the measured effect towards chance rather than away from it, so these ratios are floors. The rival explanation for the left panel is regional rather than familial recruitment, since Tunisian surnames are regionally patterned; the right panel is what separates them, because region predicts clustering in common surnames and kinship predicts it in rare ones, and the gradient runs the way kinship predicts. It does not rule region out. No ranking of surnames is drawn or computed anywhere in this build: on the common names that would head such a list the instrument is weakest, and the structural result does not need one. A body is an organisation as the register identifies it, not a ministry, so the right panel counts arrival at a directorate as well as at a ministry; a person returning to a body they already served in is not counted as their own namesake.